

KEVAL SANGHANI

Portfolio

(+91) 8160159685 ◊ Ahmedabad, Gujarat

kevalpatel0103@gmail.com ◊ [LinkedIn](#)

SKILLS

Technical	Python, Machine Learning, Deep Learning, NLP, GenAI, Agentic AI, FastAPI, SQL, LangChain, LangGraph, LangSmith, LlamaIndex, WebSocket, WebRTC, Docker, Git
Cloud	Azure (OpenAI, AI Search, AI Document Intelligence, Blob Storage, App Service), AWS (S3, RDS, OpenSearch, Bedrock)
Soft Skills	Client - Communication, Problem Solving, Logical Thinking

EXPERIENCE

AI/ML Engineer - KenexAi (formerly known as Ridgeant Technologies) November 2024 - Present
Ahmedabad, Gujarat

Agentic AI Contract Analysis Platform (Azure) – Python, LangGraph, Azure OpenAI, Azure AI Search, Azure Document Intelligence, Azure App Service, Azure Blob Storage, FastAPI, Langfuse

- Designed a **multi-agent LangGraph architecture** with an **LLM-based intent router** (structured Pydantic outputs) that classifies requests into **7 intents** and detects **prompt-injection/jailbreak** attempts.
- Built a production **RAG pipeline** on **Azure AI Document Intelligence**, **Azure OpenAI embeddings**, and **Azure AI Search (HNSW + semantic reranking)**, delivering answers with **page-level citations**.
- Automated extraction of **83 FAQ questions** and comparison of **73 contract provisions** using **YAML-driven catalogs**, letting business users change requirements without code changes.
- Developed an **LLM-powered Gap Analysis engine** that scores clauses against company-standard positions by **risk and criticality**, and generates mitigations and executive summaries.
- Engineered an **async, fault-tolerant job orchestration layer** with **SHA-256 idempotency keys**, **Azure Blob Storage** state, result caching, parallel processing, and retry/backoff for Azure rate limits.
- Deployed as a **FastAPI service on Azure App Service** (Gunicorn/Uvicorn) supporting multi-file uploads, **Apache Impala** data-lake lookups, **Langfuse** tracing, and structured logging with credential redaction.

AI Mock Interview Platform: Real-Time Voice Interview Agent – Python, Pipecat, OpenAI GPT-4o-mini, Deepgram, Cartesia, Qdrant, WebRTC, WebSockets, FastAPI, React, TypeScript

- Built an end-to-end **voice AI interviewer** on **Pipecat** with a streaming pipeline of **Deepgram STT**, **OpenAI GPT-4o-mini**, and **Cartesia TTS** for natural, spoken mock interviews.
- Implemented **RAG-based question generation**: job descriptions embedded with **OpenAI embeddings** in **Qdrant**, retrieving the **top-5 similar questions** to tailor interviews by company, role, and JD, with a fallback question set on failure.
- Set up low-latency, two-way audio over **WebRTC** (SDP offer/answer, ICE/STUN) between the browser and the bot server.
- Tuned turn-taking with **Silero VAD** and a local **Smart Turn v3 (ONNX)** model, so the bot waits for candidates to finish instead of interrupting mid-sentence.
- Built a **real-time transcript system** using **FastAPI WebSockets** with auto-reconnect, merging bot and client-side WebRTC transcripts into a live conversation view.
- Designed a **FastAPI backend** managing interview lifecycles (session creation, per-interview bot subprocesses, status) with a **React + TypeScript** frontend (Vite, Tailwind).

Explainable Text-to-SQL & RAG System (AWS) – Python, LangChain, LlamaIndex, RAG, AWS Services (S3, OpenSearch, Bedrock, RDS)

- Built a **hybrid Text-to-SQL and RAG pipeline** over a **drug-related dataset** of **102 structured tables (AWS RDS)** and **unstructured documents (AWS S3)**.
- Used **LangChain** to convert **natural language into SQL** with **LLM-based reranking**, and **AWS OpenSearch** as the **vector store** for schema and document embeddings.
- Implemented **LlamaIndex-based RAG** for contextual retrieval over **S3 documents**.
- Displayed **schema references** for SQL queries and **document names with page numbers** for unstructured data, ensuring **explainability** of results.
- **Deployed and integrated** the complete system within the **AWS ecosystem** for **scalable, production-ready usage**.

Intelligent Email Folder Recommendation System – Python, Qdrant, FAISS, Milvus, Pinecone, pgvector, Docker, FastAPI, Sentence Transformers, HNSW, IVFF

- Built an end-to-end **semantic email classification system** that recommends the **top-5 most relevant folders** for incoming emails using vector similarity search over historical email metadata.
- Embedded email subject, body, CC, and metadata into vectors stored in **Qdrant**, enabling low-latency, real-time similarity search.
- Designed **ACL-aware retrieval logic** so recommendations strictly complied with user access controls and enterprise security requirements.
- Benchmarked vector databases (**Qdrant, Milvus, FAISS, Pinecone, pgvector**) and indexing strategies (**HNSW, IVFF**), optimizing for **latency, recall, and scalability**.

Machine Learning Intern - Fxis AI
Ahmedabad, Gujarat

July 2024 – Oct 2024

AI Lawyer Ben – Legal Chatbot & Document Generation System – Python, RAG, Qdrant, Generative AI, NLP, Docker

- Contributed to a **live AI-powered legal assistant** enabling conversational querying and **automated legal document generation**.
- Designed and implemented a **RAG pipeline** to retrieve relevant legal content from embedded documents.
- Built ingestion workflows to **chunk files, generate embeddings, and store vectors** in **Qdrant** for efficient semantic search.
- Used **Docker-based services** to manage vector storage, retrieval pipelines, and deployment.
- Applied **ML, NLP, and Generative AI techniques** to improve response accuracy and system reliability.

EDUCATION

Bachelor of Computer Engineering
SAL Engineering and Technical Institute
CPI: 8.0

2021 – 2025

ACHIEVEMENTS

- Earned a **Master Certificate in Data Science** from **PW Skills**, demonstrating strong proficiency in data analysis, machine learning, and applied AI concepts.
- Certified **Databricks Generative AI Associate**, validating expertise in building and deploying **Generative AI solutions** on the Databricks platform.